



GPT-2 Fine-Tuning Dashboard

Config: lr=5e-05, batch=4, epochs=20 | Best Epoch: 18 (val_loss=0.0839) | Best Metrics: ppl=1.09, acc=73.20% | Total Time: 0h 15m


Loss Curves




Perplexity (lower is better)




Accuracy



Learning Rate Schedule



Gradient Norms



Training Time per Epoch

